

A Goal-Driven Tree-Structured Neural Model for Math Word Problems

IJCAI 2019

Zhipeng Xie^{*†}, Shichao Sun[†]

**Shanghai Key Laboratory of Data Science, Fudan University
School of Computer Science, Fudan University**

Background

- **Math word problems (MWP)**

- Automatically answer a mathematical query according to the text description
- Typical MWP
 - Short narrative that describes a partial state of the world and poses a question about an unknown quantity

A typical math word problem

Problem: Robin was making baggies of cookies with 6 cookies in each bag. If she had 23 chocolate cookies and 25 oatmeal cookies, how many baggies could she make?

Solution Expression: $(23 + 25) \div 6$ **Solution:** 8

Background

- **Seq2Seq-based solver**

- Pros

- The power of generating new expressions that do not exist in the training dataset
 - The Seq2Seq-based models exists in that they do not rely on hand-crafted features

- Cons

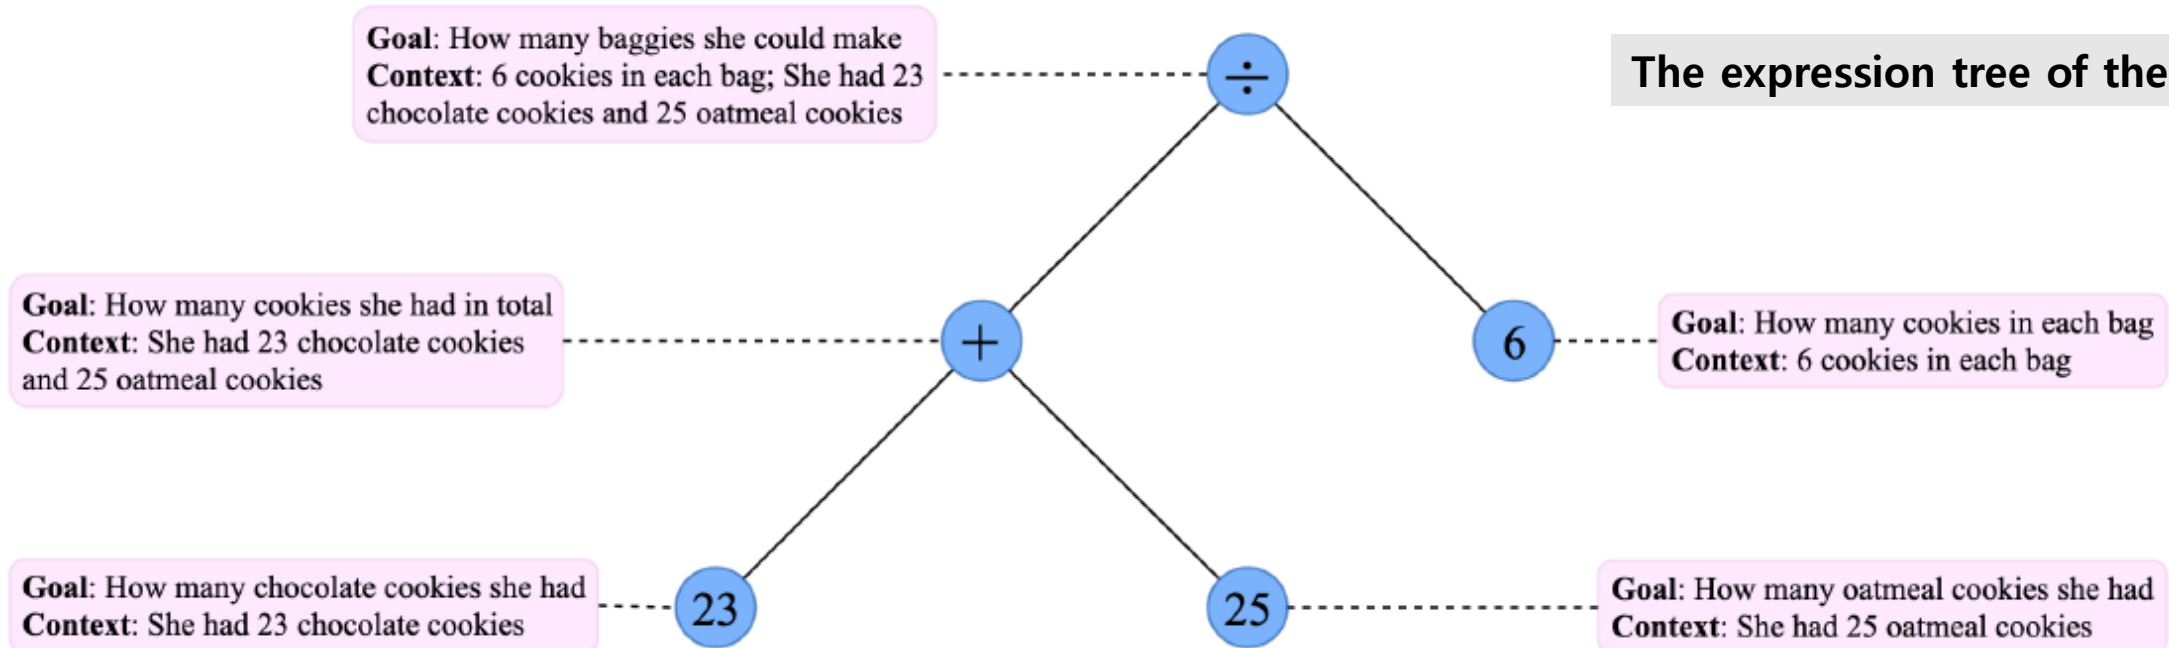
- Do not match the goal-driven mechanism in human problem solving
 - To model the tree-structured relationship of expression tree through its post order traversal sequence during decoding

Part 1. Background

- **Seq2Seq-based solver**

- Human Problem Solving

- Decompose the goal into two sub-goals combined by an operator based on the relevant information
 - Decompose the goal recursively for solving a math word problem and finally generates an expression tree



Background

- **Goal-driven mechanism in human problem solving**

- Design a novel model to generate expression tree
- Process
 - Model firstly initializes the root goal vector which represents the final goal of the problem
 - Summarizes relevant information of the problem into the context vector
 - A token is predicted using the goal vector and its context vector, which implicitly decides whether the goal should be decomposed further
 - Prediction and the goal decomposition process are repeated for them
- For a commutative operator such as "+" or "x"
 - Its right sub-goal may be the same as the left one
 - Due to its commutative property
- To address this issue
 - Our model completes the construction of the left subtree before generating the right sub-goal
 - The generation of right sub-goal takes the information of its left sibling subtree into consideration, which is encoded as a subtree embedding by a recursive neural network

- **A Goal-Driven Tree-Structured Neural Model for Math Word Problems**
 - Neural model to generate an expression tree in a human-like goal-driven way for solving math word problems
 - The first tree-structured neural model for MWPs
 - The information explicitly flows through the expression tree
 - Top-down (goal decomposition) manners
 - Bottom-up (subtree embedding) manners
 - Experimental results
 - Significantly outperforms several state-of-the-art systems on the dataset Math23K

- **Related Work**

- Rule-based methods
- Statistical machine learning methods
- Semantic parsing methods
- Deep learning methods
 - Seq2seq model with Recurrent Neural Network (RNN) in its encoder and decoder
 - Convolutional Neural Network (CNN) instead of RNN
- Huang et al. [2018]
 - The Seq2Seq model may generate spurious numbers or predict numbers at wrong positions
 - Copy-and-alignment mechanism to the standard Seq2Seq model to solve these issues
- Wang et al. [2018a]
 - Seq2Seq model always suffers from an equation duplication problem: a MWP can be solved by multiple expressions
 - An equation normalization method to solve this problem

Part 3. Problem Statement

- **Problem text P**
 - A sequence of word tokens and numeric values
 - Usually begins by describing a partial quantitative state of a world, followed by simple updates
 - Ends with a query about an unknown quantity
- **The ordered list of numeric values in P according to their order in the problem text n_P**
 - At a preprocessing step
 - All the number tokens are treated as a special word token NUM
 - Usually do not care about their exact values in solving a math word problem

Part 3. Problem Statement

- **Solution expression tree T**
 - Mathematical expression tree
 - Can be easily transformed from the solution expression
 - Capture the relations among these numeric values which are described or implied literally by the problem text
 - T may contain constant quantities, mathematical operators, and numeric values in n_P from problem text P

Part 3. Problem Statement

- **The set of mathematical operators** V_{op}
- **The set of constant quantities** V_{con}
 - Special numeric values that may occur in the solution but not in the problem text
- **The target vocabulary of** P
 - Two identical numbers since the number occurs in two different positions of P
 - Choose the occurrence of higher probability (Equation (8)) as the target

$$\text{prob}(y|\mathbf{q}, \mathbf{c}, P) = \frac{\exp(s(y|\mathbf{q}, \mathbf{c}, P))}{\sum_i \exp(s(y_i|\mathbf{q}, \mathbf{c}, P))}$$

Model Description


- **Each node \mathbf{n} in the expression tree T**
 - The goal vector \mathbf{q}
 - Instruct how the subtree root from node \mathbf{n} should be constructed
 - The subtree is generated to realize the goal
 - First predicts the token \hat{y} according to the goal vector \mathbf{q}
 - Predicted token naturally decides whether the goal should be decomposed further
 - If the predicted token is a mathematical operator, the goal will be decomposed into two sub-goals (a left sub-goal \mathbf{q}_l and a right one \mathbf{q}_r)
 - The left (right) sub-goal serves to drive the construction of the left (right) subtree of \mathbf{n}
 - The goal will be simply realized by the predicted numeric value or constant quantity
 - Such a goal decomposition process is conducted recursively just like a depth-first traversal

Part 4. Model Description

• The goal decomposition process

- The left sub-goal is generated according to the goal vector and the predicted token of its parent node

Goal-driven Tree-structured Model



Left Sub-Goal Generation

$$o_l = \sigma(\mathbf{W}_{ol} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$C_l = \tanh(\mathbf{W}_{cl} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$\mathbf{h}_l = o_l \odot C_l$$

$$g_l = \sigma(\mathbf{W}_{gl} \mathbf{h}_l)$$

$$Q_{le} = \tanh(\mathbf{W}_{le} \mathbf{h}_l)$$


$$\mathbf{q}_l = g_l \odot Q_{le}$$

Part 4. Model Description

• The goal decomposition process

- Take the information of its left sibling subtree into consideration, in addition to the parent goal and the left sub-goal

Goal-driven Tree-structured Model



Right Sub-Goal Generation

$$o_r = \sigma(\mathbf{W}_{or} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$C_r = \tanh(\mathbf{W}_{cr} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$\mathbf{h}_r = o_r \odot C_r$$

$$g_r = \sigma(\mathbf{W}_{gr} [\mathbf{h}_r, \mathbf{t}_l])$$

$$Q_{re} = \tanh(\mathbf{W}_{re} [\mathbf{h}_r, \mathbf{t}_l])$$

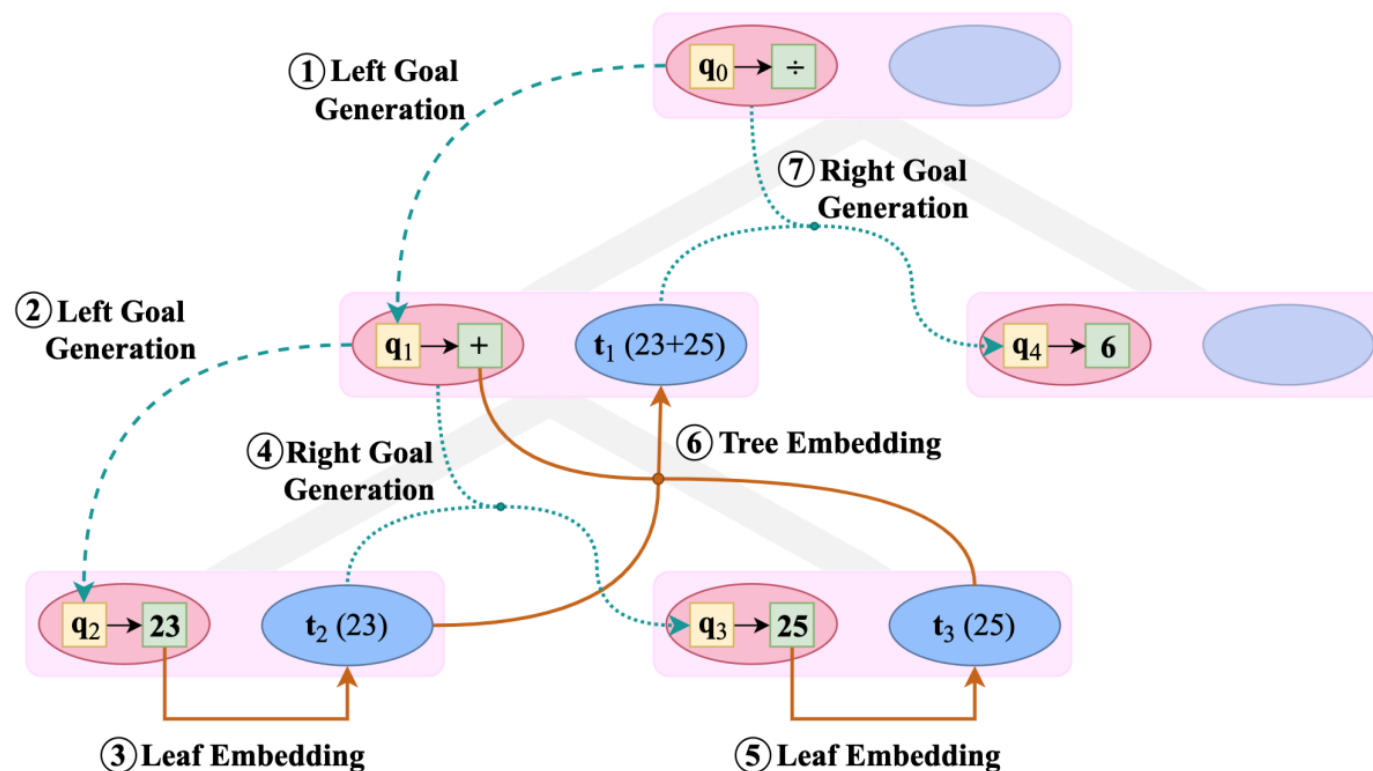
$$\mathbf{q}_r = g_r \odot Q_{re}$$

Part 4. Model Description

• The goal decomposition process

- Encode the subtree information of a non-leaf node, bottom-up RNN is defined which fuses the token embedding of its mathematical token, and the embeddings of its left and right subtrees

Goal-driven Tree-structured Model



Subtree Embedding via RNN

$$\mathbf{t} = \begin{cases} \text{comb}(\mathbf{t}_l, \mathbf{t}_r, \hat{y}) & \text{if } \hat{y} \in V_{op} \\ \mathbf{e}(\hat{y}|P) & \text{if } \hat{y} \in n_P \cup V_{con} \end{cases}$$

$$\text{comb}(\mathbf{t}_l, \mathbf{t}_r, \hat{y}) = g_t \odot C_t$$

$$g_t = \sigma(\mathbf{W}_{gt} [\mathbf{t}_l, \mathbf{t}_r, \mathbf{e}(\hat{y}|P)])$$

$$C_t = \tanh(\mathbf{W}_{ct} [\mathbf{t}_l, \mathbf{t}_r, \mathbf{e}(\hat{y}|P)])$$

Part 4. Model Description

- **Encoder**

- An input problem text $P = x_1 x_2 \dots x_n$
- Each word token x_i is firstly transformed into the corresponding word embedding \mathbf{x}_i
 - Look up an encoder embedding matrix \mathbf{M}_{sen}
- The sequence of embeddings is inputted to the Gated Recurrent Unit (GRU) [Cho et al., 2014]
- The function of a two-layer GRU
 - Produce a sequences of encoder hidden states one-by-one
 - The final hidden state \mathbf{h}_s^p has incorporated contextual information of the source token

$$\overrightarrow{\mathbf{h}}_s^p = \text{GRU}(\overrightarrow{\mathbf{h}}_{s-1}^p, \mathbf{x}_s) \quad \overleftarrow{\mathbf{h}}_s^p = \text{GRU}(\overleftarrow{\mathbf{h}}_{s+1}^p, \mathbf{x}_s) \quad \mathbf{h}_s^p = \overrightarrow{\mathbf{h}}_s^p + \overleftarrow{\mathbf{h}}_s^p$$

Part 4. Model Description

- **Root Goal Initialization and Token Embedding**

- The top-down goal decomposition process
 - Initialize the goal vector \mathbf{q}_0 of the root node \mathbf{n}_0
 - According to the hidden states of the encoder of P
 - The final hidden states of forward/backward sequence $\overrightarrow{\mathbf{h}}_n^p$ $\overleftarrow{\mathbf{h}}_0^p$

$$\mathbf{q}_0 = \overrightarrow{\mathbf{h}}_n^p + \overleftarrow{\mathbf{h}}_0^p$$

- Token embedding

$$\mathbf{e}(y|P) = \begin{cases} \mathbf{M}_{op}(y) & \text{if } y \in V_{op} \\ \mathbf{M}_{con}(y) & \text{if } y \in V_{con} \\ \mathbf{h}_{loc(y,P)}^p & \text{if } y \in n_P \end{cases}$$

Part 4. Model Description

- **Top-down Goal Decomposition**

- Context vector \mathbf{c}
 - Given a goal vector \mathbf{q} , It summarizes relevant information of the problem at hand, which is expected to help predict the token and make the following decisions
 - Weighted representation of the source tokens by a vector \mathbf{a} of attention weights a_s

$$\mathbf{c} = \sum_s a_s \mathbf{h}_s^p \qquad a_s = \frac{\exp(\text{score}(\mathbf{q}, \mathbf{h}_s^p))}{\sum_i \exp(\text{score}(\mathbf{q}, \mathbf{h}_i^p))}$$

$$\text{score}(\mathbf{q}, \mathbf{h}_s^p) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a [\mathbf{q}, \mathbf{h}_s^p]) \qquad \circ \text{ Trainable parameters } \mathbf{v}_a \mathbf{W}_a$$

- Unnormalized log probability of generating a token y from the target vocabulary \mathcal{V}^{dec}

$$s(y|\mathbf{q}, \mathbf{c}, P) = \mathbf{w}_n^\top \tanh(\mathbf{W}_s [\mathbf{q}, \mathbf{c}, \mathbf{e}(y|P)])$$

- Trainable vector \mathbf{w}_n
- Trainable matrix \mathbf{W}_s
- Token embedding of y $\mathbf{e}(y|P)$

Part 4. Model Description

- **Top-down Goal Decomposition**

- The normalization of $s(y|\mathbf{q}, \mathbf{c}, P)$ through softmax over target vocabulary

$$\text{prob}(y|\mathbf{q}, \mathbf{c}, P) = \frac{\exp(s(y|\mathbf{q}, \mathbf{c}, P))}{\sum_i \exp(s(y_i|\mathbf{q}, \mathbf{c}, P))}$$

- The predicted token implies a decision about how to realize the goal
 - If \hat{y} is a numeric value or a constant quantity, the goal is realized directly by \hat{y}
 - Otherwise (i.e., \hat{y} is an operator), the goal will be decomposed into two sub-goals

$$\hat{y} = \arg \max_{y \in V^{dec}} \text{prob}(y|\mathbf{q}, \mathbf{c}, P)$$

Part 4. Model Description

- **Left Sub-Goal Generation**

- Predicted token \hat{y} is an operator
 - The current goal \mathbf{q} will be realized by a left(right) sub-goal \mathbf{q}_l \mathbf{q}_r
- Left sub-goal is calculated by a two-layer feedforward neural network with gating mechanism

$$o_l = \sigma(\mathbf{W}_{ol} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$C_l = \tanh(\mathbf{W}_{cl} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$\mathbf{h}_l = o_l \odot C_l$$

$$g_l = \sigma(\mathbf{W}_{gl} \mathbf{h}_l)$$

$$Q_{le} = \tanh(\mathbf{W}_{le} \mathbf{h}_l)$$

$$\mathbf{q}_l = g_l \odot Q_{le}$$

- Trainable matrices \mathbf{W}_{ol} \mathbf{W}_{cl} \mathbf{W}_{gl} \mathbf{W}_{le}

- The hidden state \mathbf{h}_l
 - Parent node delivers to its left child

Part 4. Model Description

- **Right Sub-Goal Generation**

- The right sub-goal quantity of right child takes into account the left child subtree
 - The left child subtree has been generated prior to the right child owing to the essence of pre-order traversal
 - The left subtree is encoded bottom up as \mathbf{t}_l according to the recursive neural network

$$o_r = \sigma (\mathbf{W}_{or} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$C_r = \tanh (\mathbf{W}_{cr} [\mathbf{q}, \mathbf{c}, \mathbf{e}(\hat{y}|P)])$$

$$\mathbf{h}_r = o_r \odot C_r$$

$$g_r = \sigma (\mathbf{W}_{gr} [\mathbf{h}_r, \mathbf{t}_l])$$

$$Q_{re} = \tanh (\mathbf{W}_{re} [\mathbf{h}_r, \mathbf{t}_l])$$

$$\mathbf{q}_r = g_r \odot Q_{re}$$

- Trainable matrices \mathbf{W}_{or} \mathbf{W}_{cr} \mathbf{W}_{gr} \mathbf{W}_{re}
- The tree embedding of its sibling node \mathbf{t}_l

Part 4. Model Description

• Subtree Embedding via Recursive Neural Network

- A recursive neural network to encode a subtree in a bottom-up manner
- The embedding \mathbf{t} of t is defined recursively as:

$$\mathbf{t} = \begin{cases} \text{comb}(\mathbf{t}_l, \mathbf{t}_r, \hat{y}) & \text{if } \hat{y} \in V_{op} \\ \mathbf{e}(\hat{y}|P) & \text{if } \hat{y} \in n_P \cup V_{con} \end{cases} \quad \circ \text{ Subtree at hand } t$$

Case 1 : If the predicted token is an operator

- $(\hat{y} \in V_{op})$ means the subtree t must have two child subtrees t_l t_r
- The embedding of t needs to fuse the information from the operator \hat{y} , the child
- As done by the function $\text{comb}(\mathbf{t}_l, \mathbf{t}_r, \hat{y})$ with gating mechanism

$$\text{comb}(\mathbf{t}_l, \mathbf{t}_r, \hat{y}) = g_t \odot C_t$$

$$g_t = \sigma(\mathbf{W}_{gt} [\mathbf{t}_l, \mathbf{t}_r, \mathbf{e}(\hat{y}|P)])$$

$$C_t = \tanh(\mathbf{W}_{ct} [\mathbf{t}_l, \mathbf{t}_r, \mathbf{e}(\hat{y}|P)])$$

- Trainable parameter matrices \mathbf{W}_{ct} \mathbf{W}_{gt}
- The embedding of the operator $\mathbf{e}(\hat{y}|P)$

Part 4. Model Description

- **Subtree Embedding via Recursive Neural Network**

Case 2 : If \hat{y} is a numeric value or a constant quantity

- The recursion stops and t becomes a leaf node
- The embedding of subtree t is simply set as the corresponding token embedding $\mathbf{e}(\hat{y}|P)$

$$\mathbf{e}(y|P) = \begin{cases} \mathbf{M}_{op}(y) & \text{if } y \in V_{op} \\ \mathbf{M}_{con}(y) & \text{if } y \in V_{con} \\ \mathbf{h}_{loc(y,P)}^p & \text{if } y \in n_P \end{cases}$$

Part 4. Model Description

- **Training Objective**

- Minimize is the negative log-likelihood of $\mathbb{D} = \{(P^i, T^i) : 1 \leq i \leq N\}$

$$J = \sum_{(P,T) \in \mathbb{D}} -\log p(T|P)$$

- Training stage

- The decoder generates one target token at each step, in the pre-order traversal of T
- It ensures that the ground truth is used as the tree structure during training
- The conditional probability

$$p(T|P) = \prod_{t=1}^m \text{prob}(y_t | \mathbf{q}_t, \mathbf{c}_t, P)$$

$$\text{prob}(\cdot | \cdot) = \frac{\exp(s(y | \mathbf{q}, \mathbf{c}, P))}{\sum_i \exp(s(y_i | \mathbf{q}, \mathbf{c}, P))}$$

- m denotes the size of T
- The goal vector & its context vector at the t -th node \mathbf{q}_t \mathbf{c}_t

Part 5. Experiment

- **Datasets: Math23K (Wang et al., 2017)**

- 23,161 math word problems annotated with solution expressions and answers
- Crawl over 60,000 Chinese math word problems from a couple of education web sites
- Real math word problems for elementary school students
- Problems in this dataset can be solved by one linear algebra expression
- The solution expression can be easily transformed into the corresponding expression tree

Problem: Dan have 5 pens and 3 pencils, Jessica have 4 more pens and 2 less pencils than him. How many pens and pencils do Jessica have in total?

Equation: $x = 5 + 4 + 3 - 2$

Solution: 10

Problem Formulation

- A problem P can be solved by a mathematical equation E_p formed by V_p and mathematical operators
- To decrease the diversity of equations
 - Map each equation to an equation template T_p through a number mapping M_p

$$M_p \quad M : \{n_1 = 5; \quad n_2 = 3; \quad n_3 = 4; \quad n_4 = 2; \}$$

$$T_p \quad x = n_1 + n_3 + n_2 - n_4$$

Part 5. Experiment

- **Models for Comparison**

- Hybrid model w/ SNI [Wang et al., 2017]
 - Combine retrieval model & seq2seq model with significant number identification(SNI)
- Ensemble model w/ EN [Wang et al., 2018a]
 - Select the result according to models' generation probability among BiLSTM, ConvS2S, Transformer with equation normalization(EN)
- Goal-driven tree-structured MWP solver (GTS)
 - The method proposed in this paper
- GTS model w/o Subtree Embedding
 - To make clear the effect of subtree embedding component on the performance of GTS
 - Right sub-goal generation in the same way as the left sub-goal (that is, the subtree embedding component gets removed)

Experiment

- **Implementation Details**

- All the words with less than 5 occurrences are converted to a universal token "<unk>"
- The dimensionality of word embedding layer is set to 128, all hidden states for the other layers are set to 512
- Our model is trained for 80 epochs by Adam optimization algorithm [Kingma and Ba, 2014] where the mini-batch size is set to 64
- The initial value of learning rate is set to 0.001, and the learning rate will be halved every 20 epochs
- Set the dropout probability [Hinton et al., 2012] as 0.5 and weight decay as $1e - 5$ to prevent overfitting
- Last but not least, we set the beam size to 5 in beam search to generate expression trees

Part 5. Experiment

• Results and Analyses

- Answer Accuracy
 - The predicted expression tree of “ $n_0 + n_1$ ” is different from the target expression tree of “ $n_1 + n_0$ ”, but they are equivalent and their calculated values are equal
- Goal-driven mechanism is feasible for solving the math word problems
- The subtree embedding module is helpful and complementary to top-down goal decomposition process

Model comparison on answer accuracy

Model	Accuracy(%)
Hybrid model w/ SNI [Wang <i>et al.</i> , 2017]	64.7
Ensemble model w/ EN [Wang <i>et al.</i> , 2018a]	68.4
GTS model w/o Subtree Embedding	70.0
GTS model	74.3

Part 5. Experiment


- **Results and Analyses**

- Answer Accuracy vs. Training Set Size

- Check how the performance of our model varies with respect to different numbers of training instances

- Baseline

- a Seq2Seq model with attention mechanism, which contains the same encoder as GTS model but takes GRU as decoder



Part 5. Experiment

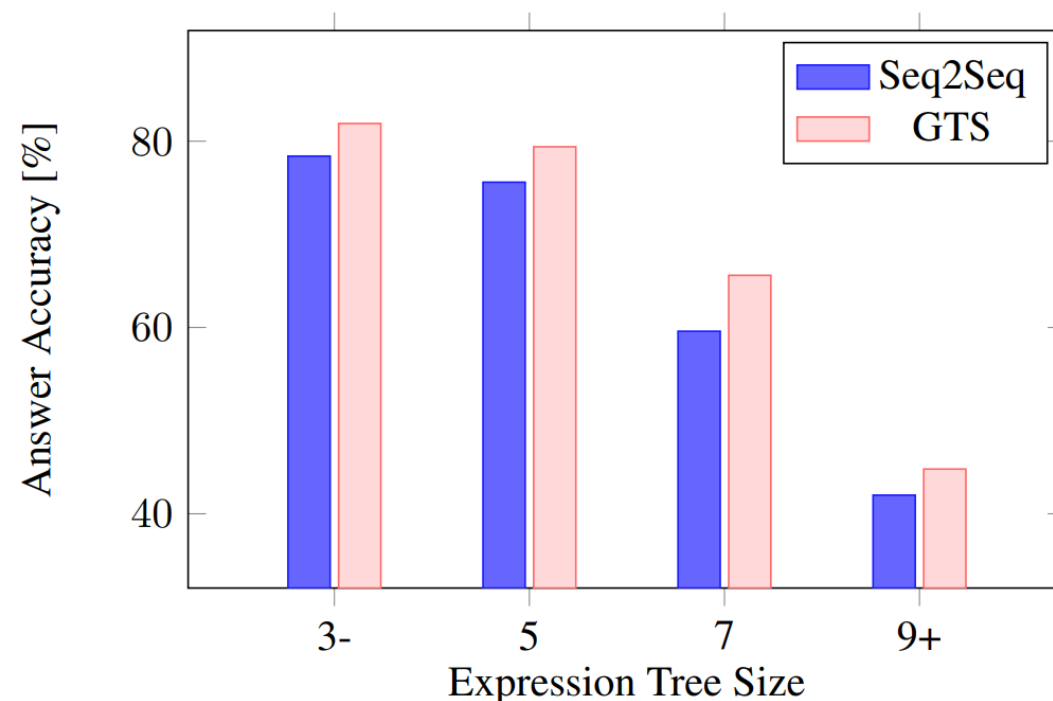
• Performance on Expression Length

- A clear tendency for answer accuracy to degrade with the growth of the problem complexity measured as the size of expression tree
- GTS model can better model the mathematical relationships of the problem in an explicit tree structure

Numbers of test instances over expression tree sizes

Expr. Tree Size	3-	5	7	9+
# Test Instances	907	2303	921	501

- Expression tree size is no less than 9, the tree will contain at least 4 mathematical operators



Part 5. Experiment

• Results on a Small Dataset

- The performance of the GTS model on another small dataset called AllArith
- Use the same 5-fold cross validation
- To improve the reproducibility, we repeat the experiment 20 times
- Due to the dataset is small, perform McNemar's test and get p-value 0.001
 - It rejects the null hypothesis and this increase is statistically significant

McNemar's test

$$H_0 : p_b = p_c \quad H_1 : p_b \neq p_c$$

$$\chi^2 = \frac{(b - c)^2}{b + c}$$

Seq2seq	GTS			
		GTS 1	GTS 2	..
	Seq2seq 1	a	b	..
	Seq2seq 2	c	d	..

Experiment

• Case Study

- Case 1: Avoid generating mathematically invalid expressions
 - Generate the tree directly, and its sequence of pre-order traversal can be guaranteed to be computable
- Case 2: Avoid predicting spurious numbers
 - Effective size of target vocabulary is set dynamically according to the specific problem
- Case 3: The subtree embedding component
 - The subtree embedding component can prevent generating the same subtree as its left sibling when the parent node is “+” or “×”

Typical cases

Case 1: The store shipped in a batch of leather shoes. $\text{NUM}(n_0 [\frac{1}{3}])$ of the total was sold on the first day, and $\text{NUM}(n_1 [\frac{3}{5}])$ of the first day's sale was sold on the second day. There were $\text{NUM}(n_2 [280])$ pairs left. How many pairs of leather shoes did the store bring in?

Seq2Seq: $\div n_2 - 1n_0 * n_0n_1$;(**error**)

GTS: $\div n_2 - -1n_0 * n_0n_1$;(**correct**)

Case 2: Of the $\text{NUM}(n_0 [697])$ combined shipment equipments of Shenzhou $\text{NUM}(n_1 [7])$ spacecraft, $\text{NUM}(n_2 [346])$ are followed, $\text{NUM}(n_3 [237])$ are updated, and the rest are newly developed. How many new equipments are there?

Seq2Seq: $- - n_0n_3n_4$;(**error**)

GTS: $- - n_0n_2n_3$;(**correct**)

Case 3: Guangming Primary School spent $\text{NUM}(n_0 [288])$ yuan on $\text{NUM}(n_1 [12])$ chairs. And then $\text{NUM}(n_2 [36])$ chairs of the same kind were bought. How much did the school spend on chairs?

GTS w/o Subtree Embedding: $\times \div n_0n_1 \div n_0n_1$;(**error**)

GTS: $\times \div n_0n_1 + n_1n_2$;(**correct**)

Conclusion

- **A Goal-Driven Tree-Structured Neural Model for Math Word Problems**
 - Motivation
 - The goal-driven mechanism in human problem solving
 - a novel neural model (called GTS) for math word problems
 - Directly predicting an expression tree
 - The information is able to flow explicitly through the expression tree by top-down goal decomposition and bottom-up subtree embedding
 - Experimental result
 - Significantly outperform previous state-of-the-art systems
 - Case study
 - Avoid generating mathematically invalid expressions and spurious numbers