

ByGPT5: End-to-End Style-conditioned Poetry Generation with Token-free Language Models

ACL 2023

Jixin Zhang, Yashar Moshfeghi

University of Strathclyde

Background

- **End-to-end fine-tuning of pre-trained language models**
 - GPT-2 or T5 on downstream tasks has been a popular training paradigm for text generation in the last few years
 - End-to-end models learn to complete a task by directly learning all steps
 - Without intermediary algorithms such as hand-crafted rules or post-processing
 - This approach has proven to be highly effective on a wide range of problems
 - Dialogue generation, Summarization, Machine translation
 - All these applications have in common
 - Only concern themselves with the generation of prosaic texts

Background

- Generating formal verse poetry remains a difficult problem
 - Strict constraints on aesthetic style such as rhyme scheme, meter and alliteration

BYGPT5 Generated quatrain with rhyme scheme, alliterations, and meter

Rhyme: 'ABAB' / Meter : "iambus(약강격)"

Alliteration: "medium"

When the sun shall set, and the moon be set,
And the night be thin and drear,
When the clouds break, and the great world be met
With the stars and moon so near,

Rhyme: 'ABBA' / Meter : "iambus"

Alliteration: "medium"

But I know the hour is near,
When the wind is in the east,
And the sun is on the feast,
And the moon is on the bier,

Rhyme: 'ABAB' / Meter : " trochee (강약격)"

Alliteration: "medium"

And the word that he spoke had been a sound of sorrow,
A sound that had stayed in the darkness without a sound,
Sound of swords and clangour of torment, call of to-morrow,
Sound of wars and thunder of battles and the blood of foes that won.

Background

- **Rhyme, Meter, Alliteration**

- Strict constraints on aesthetic style such as rhyme scheme, meter and alliteration

Rhyme & Meter

**Humpty Dumpty sat on a wall,
Humpty Dumpty had a great fall.
All the king's horses and all the king's men
Couldn't put Humpty together again.**

Alliteration

"**The daily diary of the American dream.**"
(The Wall Street Journal Commercials)

"**What we want is Watneys.**"
(Watney's Commercials)

Background

- **Rhyme, Meter, Alliteration**

- **Rhyme**

- Repetition of the same or similar sounds in the final accented syllables of words, which must be preceded by differing consonant

- **Imperfect rhymes**

- Final sounds are different or the words are identical

- **Meter**

- Rhythmic pattern within a verse
 - Accented-syllabic: Succession of stressed (—) and unstressed syllables (˘) occurs at regular intervals

Background

- **Rhyme, Meter, Alliteration**

- **Meter**

- Rhythmic pattern within a verse
 - Accented-syllabic: Succession of stressed (—) and unstressed syllables (˘) occurs at regular intervals

- The rhythmic unit: Foot

- The meter of a verse can thus be described as a sequence of feet
 - Iambic (˘ —), trochaic (— ˘), anapestic (˘ ˘ —), Dactylic (— ˘ ˘).

The sweet wild strain, the sudden start, A
Which shakes the perfumed altar's flame, B
To make its shrine a sacred name, B
And sing its praise in every heart. A

— BYGPT5

Background


- **Rhyme, Meter, Alliteration**

- **Alliteration**

- The repetition of the same consonant sounds or any vowel sounds at the beginning of words or syllables that are close together in a verse
 - Alliteration is secondary to rhyme and meter, follows less strict constraints
 - Therefore not as easily classified
 - In this work, we thus consider the level of alliteration instead, which we classify as either low, medium, or high

Background

- Generating formal verse poetry remains a difficult problem
 - Poetry generation systems rely on human guidance



CR-PO System

- TextGenRNN generate poem

	Topic	Fluency
<i>Word selection</i>		
TextGenRNN is better	17.3	24.4
IQ values are better	67.7	41.7
The two are similar	14.9	33.9
<i>Word replacement</i>		
TextGenRNN is better	14.3	3.1
CamemBERT is better	68.3	79.2
The two are similar	17.5	17.7


Table 3: Answers of human judges (%) for each method of word selection and replacement.

Background

- **Generating formal verse poetry remains a difficult problem**
 - Injecting prior knowledge in the form of hard-coded constraints to filter model outputs or modify probability distributions
 - Break the whole process down into sophisticated task-specific model pipelines

Background

- Sonnet generation framework with four distinct pipeline steps



Background

- **Cristina Garbacea and Qiaozhu Mei (2022)**
 - Incorporating prior knowledge generally becomes more difficult when heterogeneous constraints are involved or the number of constraints increases
 - Standard text-generation architectures do not lend themselves well for manually applying constraints
 - Due to the autoregressive generation of tokens from left to right, constraints at arbitrary positions cannot be implemented easily or only with additional trade-offs
 - For example, end rhymes, which come at the end of a verse, cannot be constrained in isolation due dependencies on previously generated tokens

Introduction

- **Reduce the amount of human supervision in poetry generation and explore viable end-to-end solutions**
 - Hypothesize that failing to do so far has the following root causes:
 - (i) Lack of available training data
 - Poetry corpora labeled with aesthetic styles are few and rare
 - Speculate that they do not suffice to train a generalized model
 - (ii) Unfavorable tokenization algorithms
 - Aesthetic styles of poetry such as rhyme, meter, and alliteration are often expressed at the character-level
 - Otherwise, most available off-the-shelf pre-trained models operate at the subword-level
 - ➔ Assumption:
 - Character-level models would perform similarly well at poetry generation

Introduction

- **Contributions**

- (i) Pre-train ByGPT5, to our knowledge the first decoder-only transformer for character-level language modeling
- (ii) Create QuaTrain, a large machine-labeled poetry corpus of quatrains in German and English
- (iii) Learn character-level styles better than subword-based systems, such as GPT-2 and mT5
Token-free models (ByT5, ByGPT5), while being more parameter efficient and compared to humans

Character-Level Model (ByT5)

- **ByT5: Towards a Token-Free Future with Pre-trained Byte-to-Byte Models**
 - One common choice is to assign a unique token ID to each word in a fixed finite vocabulary
 - An issue with using a fixed vocabulary of words
 - There is no obvious way to process a piece of text that contains an out-of-vocabulary word
 - E.g., map all unknown words to the same <UNK> token
 - This prevents the model from distinguishing between different out-of-vocabulary words
- ➔ A solution to the out-of-vocabulary problem: Subword tokenizers

Character-Level Model (ByT5)

- **Subword tokenizers**
 - Instead of mapping each word to a single token
 - Flexibility
 - Decompose words into smaller subword units with a goal of minimizing the total length of the token sequences for a fixed vocabulary size
 - Mispredictions
 - Typos, variants in spelling and capitalization, and morphological changes can all cause the token representation of a word or phrase to change completely
- ➔ Token-free NLP models that do not rely on a learned vocabulary to map words or subword units to tokens
 - Such models operate on raw text directly

Character-Level Model (ByT5)

- **Token-free models**

- Text data is generally stored as a sequence of bytes
 - Feeding byte sequences directly into the model enables the processing of arbitrary text sequences
- This approach is well-aligned with the philosophy of end-to-end learning
 - It endeavors to train models to directly map from raw data to predictions
- A concrete benefit in terms of model size
 - **The large vocabularies of word- or subword-level models** often result in many parameters being devoted to the vocabulary matrix
 - **A byte-level model by definition** only requires 256 embeddings

Character-Level Model (ByT5)

- **Token-free models**

- Migrating word representations out of a sparse vocabulary matrix and into dense network layers
 - Allow models to generalize more effectively across related terms (e.g., book / books) and orthographic variations
- Models with a fixed vocabulary can complicate adaptation to new languages and new terminology
- By definition, token-free models can process any text sequence


Character-Level Model (ByT5)

- **The main drawback of byte-level models**
 - Byte sequences tend to be significantly longer than token sequences
 - Computational costs of machine learning models tend to scale with sequence length
 - Much previous work on character- and byte-level models has explored ways to process long sequences efficiently using convolutions with pooling or adaptive computation time

Character-Level Model (ByT5)

- **Character-Level Model (ByT5)**

- Without a dramatically unfavorable increase in computational cost
- T5 framework : The design stays fairly close to mT5 (the multilingual variant of T5)
 - All text-based NLP problems are cast to a text-to-text format
 - Makes it simple to tackle an NLP task by generating a sequence of bytes



Character-Level Model (ByT5)

- **Changes from mT5: Dispense with the SentencePiece vocabulary**
 - Feed UTF-8 bytes directly into the model without any text preprocessing
 - The bytes are embedded to the model hidden size using a vocabulary of 256 possible byte values
 - An additional 3 IDs are reserved for special tokens: padding, end-of-sentence, and an unused <UNK> token that we include only for convention

Character-Level Model (ByT5)

- **Changes from mT5: Modify the pre-training task**
 - mT5 uses the “span corruption” pre-training objective
 - Rather than adding 100 new tokens for the sentinels, we find it sufficient to reuse the final 100 byte IDs
 - While mT5 uses an average span length of 3 subword tokens, we find that masking longer byte-spans is valuable
 - Specifically, we set our mean mask span length to 20 bytes, and show ablations of this value in Section 6

Character-Level Model (ByT5)

- **Changes from mT5: Decouple the depth of the encoder and decoder stacks**
 - Heavier encoder
 - Set our encoder depth to 3 times that of the decoder
 - Make the model more similar to encoder-only models like BERT
 - Light decoder
 - One might expect quality to deteriorate on tasks like summarization that require generation of fluent text
 - Heavy-encoder byte models performing better on both classification and generation tasks

Character-Level Model (ByT5)

- **Not all byte sequences are legal according to the UTF-8 standard**
 - Drop any illegal bytes in the model's output³ (though we never observed our models predicting illegal byte sequences in practice)
 - Like mT5, we set our sequence length to 1024 (bytes rather than tokens)
 - Train for 1 million steps over batches of 220 tokens

Character-Level Model (ByT5)

- Not all byte sequences are legal according to the UTF-8 standard
 - Drop any illegal bytes in the model’s output³ (though we never observed our models predicting illegal byte sequences in practice)

English Classification Tasks

Model	GLUE		SuperGLUE	
	mT5	ByT5	mT5	ByT5
Small	75.6	80.5	60.2	67.8
Base	83.0	85.3	72.5	74.0
Large	87.6	87.0	81.9	80.4
XL	88.7	87.9	84.7	83.2
XXL	90.7	90.1	89.2	88.6

English Generation Tasks

Model	GEM-XSum (BLEU)		TweetQA (BLEU-1)		DROP (F1 / EM)	
	mT5	ByT5	mT5	ByT5	mT5	ByT5
Small	6.9	9.1	54.4	65.7	40.0 / 38.4	66.6 / 65.1
Base	8.4	11.1	61.3	68.7	47.2 / 45.6	72.6 / 71.2
Large	10.1	11.5	67.9	70.0	58.7 / 57.3	74.4 / 73.0
XL	11.9	12.4	68.8	70.6	62.7 / 61.1	68.7 / 67.2
XXL	14.3	15.3	70.8	72.0	71.2 / 69.6	80.0 / 78.5

Models

- **End-to-end poetry generation systems for English and German by fine-tuning pretrained transformer models**
 - Two architectural variants:
 - Encoder-decoder transformers
 - Decoder-only transformers
 - Do not experiment with models with more than 400 million parameters
 - Exceed the capacity of our available GPU resources

Models

- **End-to-end poetry generation systems for English and German by fine-tuning pretrained transformer models**
 - Encoder-Decoder
 - Encoder-decoder transformers
 - Decoder-only transformers
 - Decoder-only
 - As the input for encoder-decoder models is a relatively short sequence of styles, this could lead to an underutilization of the encoder
 - Hypothesis:
Decoder-only model, with styles supplied as a prompt string, would be better suited for our task
- ➔ Refer to it as ByGPT5

Models

- **ByGPT5**
 - Modifying the architecture of ByT5 and discard its encoder component
 - Initialize the weights with the decoder of ByT53 to warm-start the training process
 - Repeat this for the three smallest variants of ByT5
 - ByT5 has an asymmetrical architecture
 - The resulting models retain only 25% of its parameters
 - Training data
 - Openwebtext2 (Gao et al., 2021) for English
 - cc100 (Conneau et al., 2020) for German

Datasets

- **Quatrain**

- Collect a range of labeled and unlabeled datasets of English and German poetry
- Label them automatically: Real quatrains with pseudo-quatrains (any consecutive sequence)

Dataset	Language	Verses	R	M	A
EPG64	English	1k	✓	✓	✗
PROSODIC	English	2k	✗	✓	✗
FORB	English	1k	✓	✓	✗
CHICAGO	English	95k	✓	✗	✗
ANTI-K	German	4k	✓	✓	✗
GRC	German	41k	✓	✗	✗
EPG	English	2.8m	✗	✗	✗
DLK	German	2.8m	✗	✗	✗
QUATRAIN	English	2.7m*	✓ †	✓ †	✓ †
	German	5.9m*	✓ †	✓ †	✓ †

* Verses may occur in multiple pseudo-quatrains.

† Labels are obtained from classifiers.

Datasets

- **Automatically labeling rhyme and meter**
 - Leverage the available labeled data and train a range of classifiers
 - Held-out gold data and subsequently use the best performing classifier for each style
 - Unlabeled poetry: English Project Gutenberg (EPG), Deutsches Lyrik Korpus (DLK) (Haider, 2021)
 - Labeled poetry: Prosodic, Chicago Rhyme Corpus (Chicago), For-better-for-verse (FORB) (Tucker, 2011), German Rhyme Corpus (GRC) (Haider and Kuhn, 2018), as well as EPG64 and Anti-K (Haider, 2021)
 - Map meters which appear less than 25 times in our labeled corpora to the special label other

Meter	Symbol
iambus	˘ –
trochee	– ˘
amphibrach	˘ – ˘
anapest	˘ ˘ –
dactyl	– ˘ ˘
alexandrine	˘ – ˘ – ˘ – ˘ – ˘ – ˘ –

Datasets

- **Automatically labeling rhyme and meter**
 - Meter classification
 - A multiclass classification problem with a single verse as input
 - Classify the meter of a quatrain by choosing the dominant meter among the verses
 - Rhyme classification
 - A binary classification problem with two verses separated by a special token as input
 - The rhyme scheme by determining which verses rhyme and which do no

F1-Score on classifying rhyme and meter

Model	Rhyme	Meter
CANINE-C	98.05	58.49
XLM-R	97.22	54.65
MBERT	97.17	49.01

Datasets

- **Labeling for alliteration**

- The quantification of the level of alliteration in a document is a long known research problem (Skinner, 1939; Leavitt, 1976; Blain, 1987; Benner, 2014)
- Let v_i be the atomic units of sound in verse v , Blain (1987) quantify alliteration $f(\cdot)$: a similarity function of two sounds; the default simply testing for equality
- $\text{allit}(\cdot)$ counts alliterative sounds in a verse, applies a distance penalty, and normalizes the score to $[0, 1]$

$$\text{allit}(v) = \frac{\sum_{i=1}^{|v|} \sum_{j=i+1}^{|v|} \frac{f(v_i, v_j)}{j-i}}{\sum_{i=1}^{|v|} \sum_{j=i+1}^{|v|} \frac{1}{j-i}}$$

Datasets

- **Labeling for alliteration**

- Consider initial phonemes of words, as well as all further stressed phonemes as atomic sound units ν_i , and to determine phonemes and stress
- Use a grapheme-to-phoneme conversion model (Zhu et al., 2022)
- Conduct an internal study to determine several intensity thresholds based on a sample of quatrains
- Classify the alliteration level of a quatrain as **low** if the score is below 0.05, **medium** if the score is below 0.1, and **high** if it is above that

Experiment

- **Automatic Evaluation**

- Rhyme schemes (AABB, ABAB, ABBA, and ABCB)
- The most popular meters per language
 - iambus, trochee, anapest, and dactyl for English
 - iambus, trochee, and alexandrine for German
- All levels of alliteration
- Create 75 poems per model for each possible combination
- To find out if styles are properly reflected in generated quatrains we reuse the classifiers

Experiment

- **Automatic Evaluation**

- Rhyme Score
 - Computes the recall of verses that should rhyme in a quatrain, as well as the recall of verses that should not and takes their arithmetic average
- Alliteration Score
 - 1 if a quatrain has the correct alliteration level, else 0.
- Meter Score
 - The fraction of verses with correctly classified meters
- Coherence
 - Uses BERT for next sentence prediction to assess discourse relations of verses
 - The score is the fraction of consecutive verse pairs that are correctly classified to come after one another

Experiment

- Automatic Evaluation





Figure 3: Automatic evaluation results for all models on English and German.

Experiment

- Human Evaluation



Model	Rhyme	Human Likeness
Human	0.84	0.89
ByGPT5	0.57	0.55
GPT-2	0.35	0.55
ByT5	0.45	0.19
mT5	0.28	0.32

Figure 4: Distributions of BWS scores for rhyme, human likeness, and meter annotations through kernel density estimation. Scores range from -1 (very bad) to 1 (very good). The “•” markers denote expected values.

Conclusion

- **Automatic Evaluation**

- End-to-end style conditioned poetry generation systems for quatrains in English and German
 - Generate poetry without the need for human supervision, except for the use of poetic training data
- Token-free decoder-only language model, and show that fine-tuning it on a custom poetry corpus outperforms other models, such as GPT-2, mT5, and ByT5
 - Perform favorably against human poets in our constrained setting
- Future Work
 - Extend our system to other poetic forms such as sonnets, limericks, or villanelles

Appendix

- Additional example poems generated with ByGPT5 (base) in German and English.

German	English
<i>Ein Reiter steht am Hafen,</i> A <i>Der schaut die Flut nicht an,</i> B <i>Er hört die Schiffer schlafen</i> A <i>Im stillen Ozean.</i> B	<i>With languid smile, the stealing tear retires,</i> A <i>And the slow fading light on trembling fires!</i> A <i>Now she receives the golden circlet round,</i> B <i>And fills the woven chambers with a sound;</i> B
<i>Schweigend stehn die Burgen nieder,</i> A <i>Und die Lüfte sind verhallt,</i> B <i>Und die Trommeln klingen wider,</i> A <i>Und die Büchsen knallen halt.</i> B	<i>The first who learned the lesson there</i> A <i>Had learned to scoff and scorn to sneer,</i> A <i>And that the learned might have been</i> B <i>A shameless woman and a queen.</i> B

Appendix

- Additional example poems generated with ByGPT5 (base) in German and English.

German	English
<p>Der Greis erbebt, die Hand erstarrt, Die Kinder schauern vor dem Sterben; Die Stimme bricht, die Thräne fällt, Sie sieht ihm nach mit nassem Beben.</p>	<p>Then came the labour of the day within, The gray beginning of the week, And down we went with hope and terror sin, And not a word to say and speak.</p>
<p>Die Strömung wiederum durch alle Glieder dringt, Und alles, was da lebt und wallt und leuchtet, singt. Da steigt ein Palmenstrauch aus dem erhabnen Lichte, Er schwimmt auf einer Fluth und singet in Gedichte.</p>	<p>For the stars are in the sky; And the stars have gone to die With their songs of joy and fear, With their music and their cheer.</p>